

PA4

October 29, 2022

```
[67]: ## import libraries
import numpy as np
import matplotlib.pyplot as plt
import skimage

#torch
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
import torch.nn.functional as F
from torchvision.datasets import MNIST

#import pca
from sklearn.decomposition import PCA
from sklearn.metrics import mean_squared_error as mse
```

```
[2]: batch_size = 256
epochs = 10
learning_rate = 0.001
```

```
[3]: #Download the dataset
train_dataset = MNIST(root='./data',
                      train=True,
                      transform=transforms.ToTensor(),
                      download=True)


test_dataset = MNIST(root='./data',
                    train=False,
                    transform=transforms.ToTensor())
```

```
[4]: # Data loader
train_loader = torch.utils.data.DataLoader(dataset=train_dataset,
                                           batch_size=batch_size,
                                           shuffle=True)

test_loader = torch.utils.data.DataLoader(dataset=test_dataset,
```

```
batch_size=batch_size,  
shuffle=False)
```

```
[5]: #code taken from the official documentation  
labels_map = {  
    0: "0",  
    1: "1",  
    2: "2",  
    3: "3",  
    4: "4",  
    5: "5",  
    6: "6",  
    7: "7",  
    8: "8",  
    9: "9",  
}  
figure = plt.figure(figsize=(8, 8))  
cols, rows = 3, 3  
for i in range(1, cols * rows + 1):  
    sample_idx = torch.randint(len(train_dataset), size=(1,)).item()  
    img, label = train_dataset[sample_idx]  
    figure.add_subplot(rows, cols, i)  
    plt.title(labels_map[label])  
    plt.axis("off")  
    plt.imshow(img.squeeze(), cmap="gray")  
plt.show()
```



```
[6]: train_images = train_dataset.data
train_labels = train_dataset.targets
print(f"Training image shape:{train_images.shape}")
print(f"Training Targets shape:{train_labels.shape}")

val_images = test_dataset.data
val_labels = test_dataset.targets
print(f"validation image shape:{val_images.shape}")
print(f"validation Targets shape:{val_labels.shape}")
```

Training image shape:torch.Size([60000, 28, 28])

```
Training Targets shape:torch.Size([60000])
validation image shape:torch.Size([10000, 28, 28])
validation Targets shape:torch.Size([10000])
```

0.1 Q1

```
[7]: ##PCA and Autoencoder
```

```
[8]: #first 30 eigenvectors corresponding to top 30 eigenvalues
pca_comps_n =30
print('PCA with', pca_comps_n , 'principal components')
train_data = np.asarray(train_images)/255
test_data = np.asarray(val_images)/255
total_data = np.concatenate((train_data,test_data))
PCA_data = total_data.reshape(-1,28*28)
# PCA_data = train_data.reshape(-1,28*28)


pca1 = PCA(n_components = pca_comps_n)
pca1.fit(PCA_data)
train_pca = pca1.transform(PCA_data)
reconstructed_data = pca1.inverse_transform(train_pca)
PCA_error = mse(PCA_data,reconstructed_data)
print('Reconstruction error made by PCA: ',PCA_error)
```

```
PCA with 30 principal components
Reconstruction error made by PCA: 0.01805640184447451
```

```
[9]: #0-5,1-0,2-4,3-1,4-9,5-2,7-3,13-6,15-7,17-8
keys = [1,3,5,7,2,0,13,15,17,4]
```

```
[10]: for i in range(1, cols * rows + 1):
        sample_idx = torch.randint(len(train_dataset), size=(1,)).item()
        img, label = train_dataset[sample_idx]
        figure.add_subplot(rows, cols, i)
        plt.title(labels_map[label])
        plt.axis("off")
        plt.imshow(img.squeeze(), cmap="gray")
plt.show()
```

4



```
[11]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```





```
[12]: from numpy import uint8





for t,j in zip(range(10),keys):
    input_image = np.asarray(255*PCA_data[j],dtype=uint8).reshape(28,28)
    reconstruct_image = np.asarray(255*reconstructed_data[j],dtype=uint8).
    ↪reshape(28,28)
    plt.subplot(1,2,1)

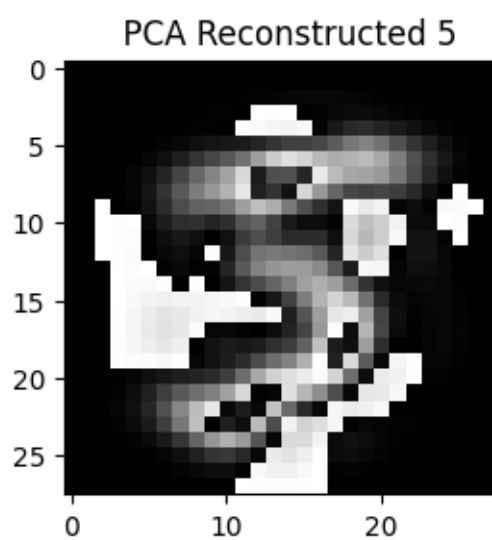



    plt.imshow(input_image,'gray')
    plt.title(' Original '+str(t))
    plt.subplot(1,2,2)





    plt.imshow(reconstruct_image,'gray')
    plt.title(' PCA Reconstructed '+str(t))
    plt.show()
```

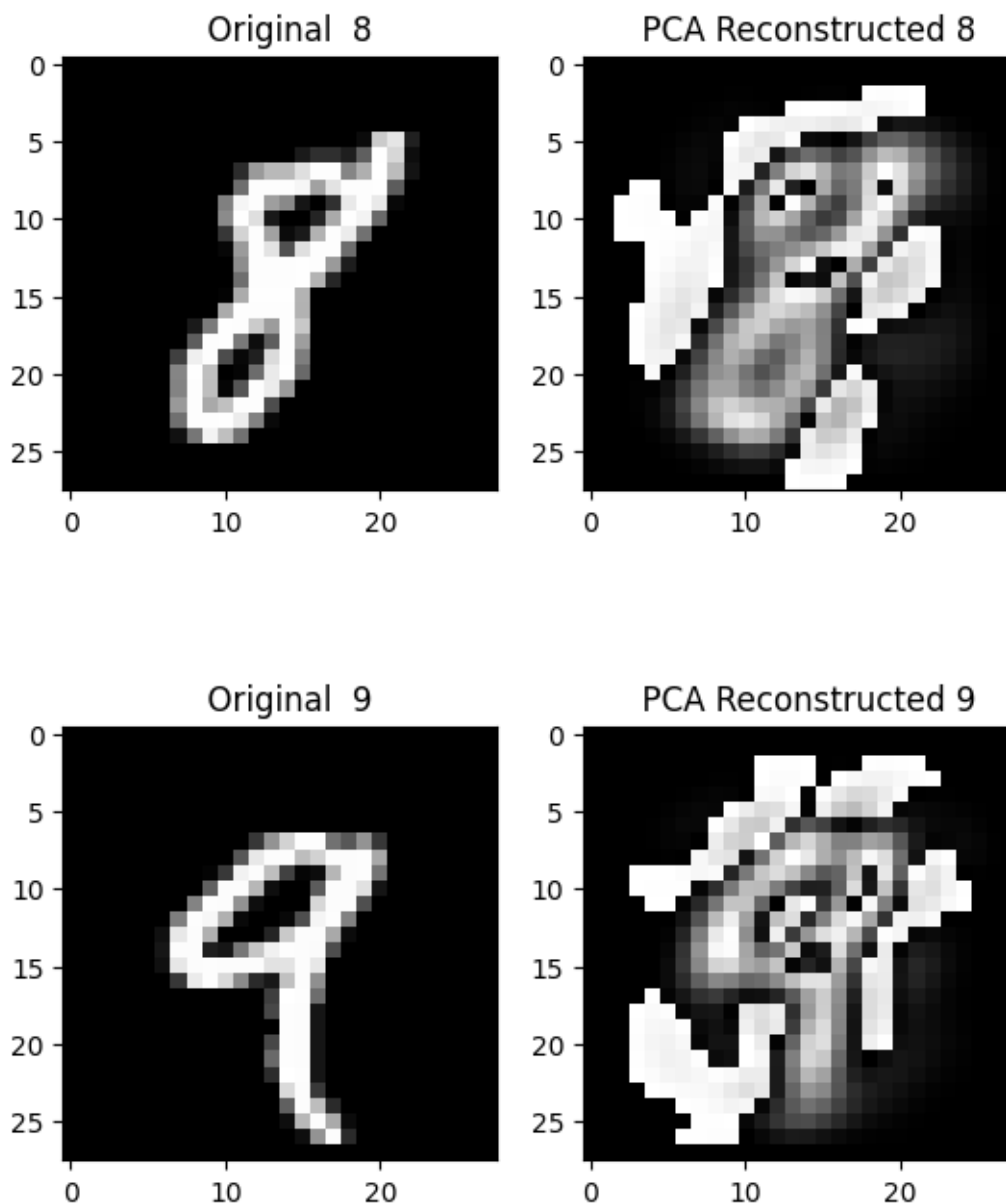
5











```
[13]: from torch import flatten
```

```
[14]: class AE1(nn.Module):  
    def __init__(self):  
        super(AE1, self).__init__()  
        self.encoder = nn.Sequential(  
            nn.Linear(784,512),  
            nn.ReLU(),  
            nn.Linear(512,256),  
            nn.ReLU(),
```

```

        nn.Linear(256,128),
        nn.ReLU(),
        nn.Linear(128,30),
        nn.ReLU())

    self.decoder =nn.Sequential(
        nn.Linear(30,128),
        nn.ReLU(),
        nn.Linear(128,256),
        nn.ReLU(),
        nn.Linear(256,784),
        nn.ReLU())

    def forward(self,x):
        x = flatten(x,1)
        encoded=self.encoder(x.float())
        reconstructed =self.decoder(encoded)
        return reconstructed , encoded

```

```

[15]: # Loss function
      lossfn = nn.MSELoss()

```

```

[16]: from torch import flatten
      def
      ↪train(model,device,train_dataloader,optimizer,lossfn,lambda_reg=0,sparse=False,denoise=False,
      ↪3, q5_flag = False):
          model.train() #setting the model in training mode
          #initializing the total training loss to 0
          train_loss = 0
          #loop over the training set
          for batch_idx, (data,label) in enumerate(train_dataloader): # (data,label):
      ↪ Training data for that batch
              if denoise==True:
                  img = data.clone()
                  data = add_noise(img,noise_val)
                  data = data.to(device)
              else:
                  (data,label) = (data.to(device),label.to(device)) #sending the
      ↪data to the device we've chosen

          reconstruction,encoded = model(data) #our reconstruction
          if q5_flag==True:
              loss = lossfn(reconstruction,data) #loss
          else:
              loss = lossfn(reconstruction,flatten(data,1)) #loss
          if sparse==True:
              loss += lambda_reg*torch.linalg.norm(encoded,1)
          optimizer.zero_grad() #zeroing out the gradients before backprop

```

```

        loss.backward()      #backprop from the loss
        optimizer.step()     #updating the weights
        train_loss += loss/len(train_dataloader)
    return train_loss #returning loss

```

```

[17]: def _
    ↪test(model,device,test_dataloader,lossfn,lambda_reg=0,sparse=False,denoise=False,noise_val=
    ↪3, q5_flag = False):
        model.eval()
        test_loss = 0
        with torch.no_grad():
            for (data,label) in test_dataloader: # (data,label): Test data for that
    ↪batch
                if denoise==True:
                    img = data.clone()
                    data = add_noise(img,noise_val)
                    data = data.to(device)
                else:
                    (data,label) = (data.to(device),label.to(device)) #sending the
    ↪data to the device we've chosen
                    #perform forward pass and compute the loss
                    reconstruction,encoded = model(data) #our prediction
                    if q5_flag==True:
                        loss = lossfn(reconstruction,data) #loss
                    else:
                        loss = lossfn(reconstruction,flatten(data,1)) #loss
                    if sparse==True:
                        loss += lambda_reg*torch.linalg.norm(encoded,1)
                    test_loss += loss/len(test_dataloader)
    return test_loss #returning loss

```

```

[18]: def _
    ↪train_test(model,device,train_loader,test_loader,optimizer,lossfn,lambda_reg=0,sparse=False,
    ↪3,q5_flag=False):
        train_losses = []
        test_losses = []
        for epoch in range(epochs+1):
            #train the model
            train_loss = _
    ↪train(model,device,train_loader,optimizer,lossfn,lambda_reg,sparse,denoise,noise_val,q5_flag)
            train_losses.append(train_loss.item())
            #test the model
            test_loss = _
    ↪test(model,device,test_loader,lossfn,lambda_reg,sparse,denoise,noise_val,q5_flag)
            test_losses.append(test_loss.item())
    return train_losses, test_losses

```

```
[19]: def plot_reconstructed_image(model,device,img,model_name,q5_flag=False):
#     img = torch.from_numpy(img)
    with torch.no_grad():
        if q5_flag == False:
            if (device==torch.device("cuda")):
                img = img.view(-1,28,28).cuda().float()
            else:
                img = img.view(-1,28,28).float()
        else:
            if (device==torch.device("cuda")):
                img = img.reshape(1,1,28,28).cuda().float()
            else:
                img = img.reshape(1,1,28,28).float()
        reconstructed_image,encoded = model.forward(img) #as it is a single
        →image we directly run the forward pass
        reconstructed_image=reconstructed_image.detach().cpu().numpy()
        img = img.reshape(28,28).detach().cpu().numpy()
        plt.subplot(1,2,1)
        plt.imshow(img , cmap='gray')
        plt.title('input image')
        plt.subplots_adjust(right=1.5)
        plt.subplot(1,2,2)
        plt.imshow(reconstructed_image.reshape(28,28),cmap = 'gray') #our
        →reconstructed image
        plt.title("Reconstructed image using "+ str(model_name))
        plt.show()
```

```
[20]: def plot_losses(train_losses, test_losses, model_name):
    train_interval = int(len(train_losses)/epochs)
    plt.plot(np.asarray(train_losses)[::train_interval])
    plt.title("MSE train loss using "+ str(model_name))
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.show()

    test_interval = int(len(test_losses)/epochs)
    plt.plot(np.asarray(test_losses)[::test_interval])
    plt.title("MSE test loss using "+ str(model_name))
    plt.xlabel("Epoch")
    plt.ylabel("loss")
    plt.show()
```



```
[21]: model_Q1 = AE1().to(device)
optimizer = torch.optim.Adam(model_Q1.parameters(), lr=learning_rate)
train_losses_AE_Q1 , test_losses_AE_Q1 =
    →train_test(model_Q1,device,train_loader,test_loader,optimizer,lossfn)
```

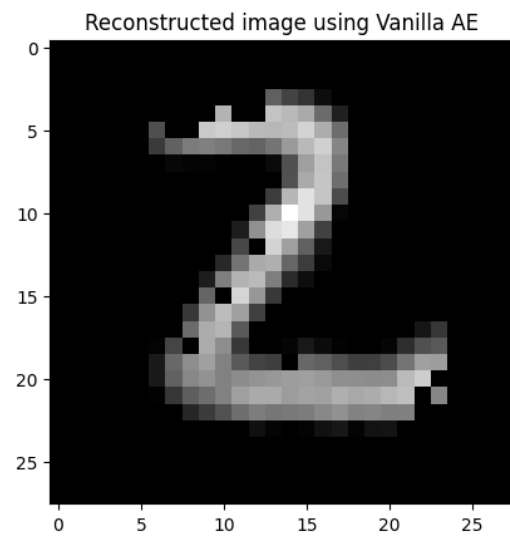

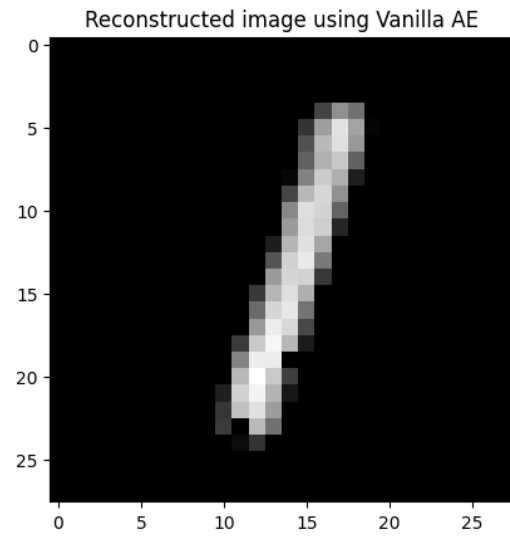
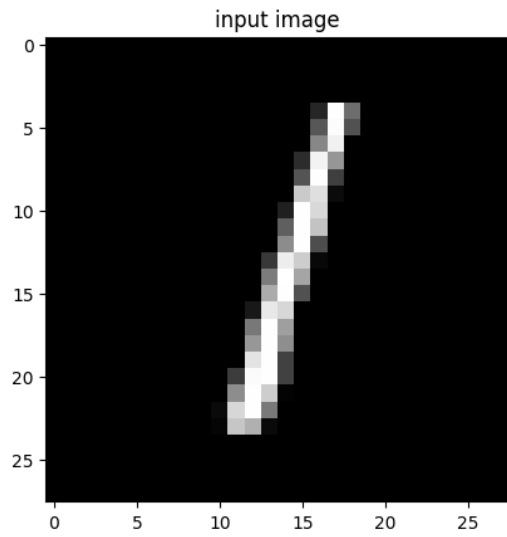
```
[22]: '''
0-7
1-2
2-1
3-0
4-4
7-9
8-5
11-6
18-3
84-8
'''
```





```
[22]: '\n0-7\n1-2\n2-1\n3-0\n4-4\n7-9\n8-5\n11-6\n18-3\n84-8\n'
```





```
[23]: keys_test=[3,2,1,18,4,8,11,0,84,7]
```





```
[24]: import random
# index = random.randint(0,9999)
for i in keys_test:
    test_image = test_loader.dataset.data[i, :, :].clone()
    plot_reconstructed_image(model_Q1,device,test_image, model_name="Vanilla_
↪AE")
# print(type(test_image))
```
















```
[25]: plot_losses(train_losses_AE_Q1, test_losses_AE_Q1, model_name = "Vanilla AE")
```





```
[26]: #MSE recomnstruction error for vanilla AE
mse_error = test(model_Q1,device,test_loader,lossfn)
print("MSE Reconstruction error for Vanilla AE is ", mse_error.item())
```

MSE Reconstruction error for Vanilla AE is 0.016686905175447464


0.2 Q2

```
[27]: class AE_Q2(nn.Module):
    def __init__(self,hidden_layer):
        super(AE_Q2, self).__init__()

        self.encoder = nn.Sequential(
            nn.Linear(784,hidden_layer),
            nn.ReLU())
        self.decoder =nn.Sequential(
            nn.Linear(hidden_layer,784),
            nn.ReLU())
```


```
def forward(self,x):
    x = flatten(x,1)
    encoded=self.encoder(x.float())
    reconstructed =self.decoder(encoded)
    return reconstructed    , encoded
```

```
[28]: from skimage import data
astronaut = data.astronaut()
plt.imshow(astronaut)
plt.show()
print(astronaut.shape)
from skimage.color import rgb2gray
grayscale_astro = rgb2gray(astronaut)
plt.imshow(grayscale_astro)
print(grayscale_astro.shape)
plt.show()
# flower = cv2.imread("flower.png")
# flower = cv2.cvtColor(flower, cv2.COLOR_BGR2GRAY)
# plt.imshow(flower)
```



(512, 512, 3)

(512, 512)



```
[29]: import skimage.transform
from skimage import img_as_ubyte
grayscale_astro = np.asarray(grayscale_astro)
non_digit_image = torch.from_numpy(img_as_ubyte(skimage.transform.
    ↳resize(grayscale_astro, (28,28))))
```


```
[30]: X = np.random.normal(loc=128,scale=10,size=(28,28)) #Initializing a 28x28
    ↳matrix with gaussian noise with mean = 128, std=10
#converting X to Tensor
if(device==torch.device("cuda")):
    noisy_image = torch.from_numpy(X).reshape(1,1,28,28).cuda().float()
else:
    noisy_image = torch.from_numpy(X).reshape(1,1,28,28).float()
```




```
[31]: x = [64,128,256] #size of the hidden layer
for hidden_layer in x:
    model_Q2 = AE_Q2(hidden_layer=hidden_layer).to(device)
    optimizer = torch.optim.Adam(model_Q2.parameters(), lr=learning_rate)
    train_losses_AE_h , test_losses_AE_h =
    ↳train_test(model_Q2,device,train_loader,test_loader,optimizer,lossfn)
```




```




    plot_losses(train_losses_AE_h, test_losses_AE_h, model_name = "AE with_
    ↪hidden_layer="+str(hidden_layer))
    for i in keys_test:
        test_image = test_loader.dataset.data[i, :, :].clone()
        plot_reconstructed_image(model_Q2,device,test_image, model_name="AE_
    ↪with hidden_layer="+str(hidden_layer))
        plot_reconstructed_image(model_Q2,device,non_digit_image, model_name="AE_
    ↪with hidden_layer="+str(hidden_layer)+"with non digit image")
        plot_reconstructed_image(model_Q2,device,noisy_image, model_name="AE with_
    ↪hidden_layer="+str(hidden_layer)+"with noisy image")





```








































































































```
[ ]:
```

0.3 Q3

```
[32]: class AE_Q3(nn.Module):
    def __init__(self):
        super(AE_Q3, self).__init__()
        self.encoder = nn.Sequential(
            nn.Linear(784,1225),
            nn.ReLU())
        self.decoder =nn.Sequential(
            nn.Linear(1225,784),
            nn.ReLU())

    def forward(self,x):
        x = flatten(x,1) #flatten the image to a 784x1 vector
        encoded_input = self.encoder(x.float())
        reconstructed_input = self.decoder(encoded_input)
        return reconstructed_input,encoded_input

[33]: def avg_hl_activations(model,test_dataloader,model_name):
    model.eval()
    avg_act_val = 0
    with torch.no_grad():
        for (data,label) in test_dataloader:
            (data,label) = (data.to(device),label.to(device))
            reconstruction,encoded = model(data) #our prediction
            avg_act_val += float(torch.mean(encoded))
    avg_act_val /= len(test_dataloader)
    print("The average activation of " + str(model_name)+" is",avg_act_val)

[34]: def encoder_decoder_filters_plots(model,model_name,device):
    with torch.no_grad():
        encoder_filters = model.encoder[0].weight.detach().cpu().numpy()
        decoder_filters = model.decoder[0].weight.detach().cpu().numpy()
        #plot the encoder and decoder weights as an image for 0th neuron
        plt.imshow(encoder_filters[0].reshape(28,28), cmap='gray')
        plt.colorbar()
        plt.title('Encoder Filters for '+str(0)+'th neuron of '+
↪str(model_name))
        plt.show()

        plt.imshow(decoder_filters[:,0].reshape(28,28), cmap='gray')
        plt.colorbar()
```

```

plt.title('Decoder Filters for '+str(0)+'th neuron of '+
↳str(model_name))
plt.show()

```

```

[35]: def visualize_activations(model, test_dataloader, model_name, device, hidden_layer):
↳ #visualize the activations
    # data_ind = np.random.randint(low=0, high=9999, size=5)
    for i, ind in enumerate(keys_test):
        test_image = test_dataloader.dataset.data[ind].clone()
        test_label = test_dataloader.dataset.targets[ind].clone()
        with torch.no_grad():
            if(device == torch.device("cuda")):
                test_image = test_image.reshape(1,1,28,28).cuda().float()
            else:
                test_image = test_image.reshape(1,1,28,28).float()
            reconstructed_image, encoded = model.forward(test_image)
            encoded = encoded.detach().cpu().numpy()
            plt.imshow(encoded.reshape(int(np.sqrt(hidden_layer)),int(np.
↳sqrt(hidden_layer))), cmap='gray')
            str_title = "Activation for digit "+str(test_label.item())
            plt.title(str_title)
            plt.show()


```

```


[36]: lambda_reg_vals=[0.000001, 0.001, 0.1]
for lambda_reg in lambda_reg_vals:
    model_Q3 = AE_Q3().to(device)
    optimizer = torch.optim.Adam(model_Q3.parameters(), lr=learning_rate)
    train_losses_AE_Q3 , test_losses_AE_Q3 =
↳train_test(model_Q3, device, train_loader, test_loader, optimizer, lossfn, lambda_reg, sparse=True)
    plot_losses(train_losses_AE_Q3, test_losses_AE_Q3, model_name =
↳"overcomplete sparse AE with sparsity parameter "+str(lambda_reg))
    for i in keys_test:
        test_image = test_loader.dataset.data[i, :, :].clone()
        plot_reconstructed_image(model_Q2, device, test_image, model_name="AE
↳with hidden_layer="+str(hidden_layer))
        plot_reconstructed_image(model_Q3, device, test_image,
↳model_name="overcomplete sparse AE with sparsity parameter "+str(lambda_reg))
        avg_hl_activations(model_Q3, test_loader, "overcomplete sparse AE with
↳sparsity parameter "+str(lambda_reg))
        encoder_decoder_filters_plots(model_Q3, "sparse AE with sparsity
↳parameter "+str(lambda_reg), device)
        visualize_activations(model_Q3, test_loader, "sparse AE with sparsity
↳parameter "+str(lambda_reg), device, 1225)





```





MSE train loss using overcomplete sparse AE with sparsity parameter 1e-06











MSE test loss using overcomplete sparse AE with sparsity parameter 1e-06





















The average activation of overcomplete sparse AE with sparsity parameter $1e-06$ is 0.27716308906674386


Encoder Filters for 0th neuron of sparse AE with sparsity parameter $1e-06$





Decoder Filters for 0th neuron of sparse AE with sparsity parameter $1e-06$





Activation for digit 0




































MSE train loss using overcomplete sparse AE with sparsity parameter 0.001










MSE test loss using overcomplete sparse AE with sparsity parameter 0.001





















The average activation of overcomplete sparse AE with sparsity parameter 0.001 is 0.004904146178159863


Encoder Filters for 0th neuron of sparse AE with sparsity parameter 0.001





Decoder Filters for 0th neuron of sparse AE with sparsity parameter 0.001





Activation for digit 0





































MSE train loss using overcomplete sparse AE with sparsity parameter 0.1











MSE test loss using overcomplete sparse AE with sparsity parameter 0.1




















The average activation of overcomplete sparse AE with sparsity parameter 0.1 is 0.0004624879096809309


Encoder Filters for 0th neuron of sparse AE with sparsity parameter 0.1





Decoder Filters for 0th neuron of sparse AE with sparsity parameter 0.1































```
[37]: encoder_decoder_filters_plots(model_Q2,"standard AE with hidden_layer size_↵  
↵256",device)
```


Encoder Filters for 0th neuron of standard AE with hidden_layer size 256





Decoder Filters for 0th neuron of standard AE with hidden_layer size 256





```
[38]: visualize_activations(model_Q2,test_loader,"Standard AE with_  
      ↪hidden_layer=256",device,256)
```




























0.4 Q4

```
[39]: def add_noise(img, noise_val):
        noise = torch.randn(img.size())*noise_val
        noisy_img = img + noise
        return noisy_img
```



```
[41]: noise_vals = [0.3,0.5,0.8,0.9]
        for noise_val in noise_vals:
            model_Q4 = AE_Q2(256).to(device)
            optimizer = torch.optim.Adam(model_Q4.parameters(), lr=learning_rate)
            train_losses_AE_Q4 , test_losses_AE_Q4 =
            ↪train_test(model_Q4,device,train_loader,test_loader,optimizer,lossfn,denoise=True,noise_val=noise_val)
            plot_losses(train_losses_AE_Q4, test_losses_AE_Q4, model_name = "Denoising_
            ↪AE with noise level "+str(noise_val))
            #MSE recomnstruction error for vanilla AE
            mse_error = test(model_Q4,device,test_loader,lossfn)
            print(" MSE for the ",noise_val,"=", mse_error.item())
            index = random.randint(0,9999)
            test_image = test_loader.dataset.data[index, :, :].clone()
```

```
test_image = add_noise(test_image,noise_val)
plot_reconstructed_image(model_Q4,device,test_image, model_name="Denoising_
↪AE with noise level "+str(noise_val))
encoder_decoder_filters_plots(model_Q3,"Denoising AE with noise level_
↪"+str(noise_val),device)
```







MSE for the 0.3 = 0.003318117465823888





Encoder Filters for 0th neuron of Denoising AE with noise level 0.3





Decoder Filters for 0th neuron of Denoising AE with noise level 0.3









MSE for the 0.5 = 0.0038982063997536898





Encoder Filters for 0th neuron of Denoising AE with noise level 0.5





Decoder Filters for 0th neuron of Denoising AE with noise level 0.5









MSE for the 0.8 = 0.006269432138651609





Encoder Filters for 0th neuron of Denoising AE with noise level 0.8





Decoder Filters for 0th neuron of Denoising AE with noise level 0.8









MSE for the 0.9 = 0.006520132068544626



Encoder Filters for 0th neuron of Denoising AE with noise level 0.9



Decoder Filters for 0th neuron of Denoising AE with noise level 0.9



```
[ ]: # Convolutional Autoencoders
```

0.5 Q5

```
[42]: class conv_AE_with_unpooling(nn.Module): #define unpooling outside the decoder
    ↪and separately in forward nn.Sequential just takes one input

    def __init__(self): #class constructor
        super(conv_AE_with_unpooling,self).__init__() #calls the parent
    ↪constructor

        #initializing the encoder module
        self.encoder_conv1 = nn.Sequential(
            nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        ) # 28x28x1 to 14x14x8
        self.encoder_conv2 = nn.Sequential(
            nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
```

```

        nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
    ) #14x14x8 to 7x7x16
    self.encoder_conv3 = nn.Sequential(
        nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
    ) #7x7x16 to 3x3x16

    #initializing the decoder module
    self.decoder_conv1 = nn.Sequential(nn.Identity()) #7x7x16 to 7x7x16
    self.decoder_conv2 = nn.Sequential(
        nn.Conv2d(16,8, kernel_size = 3, stride = 1,padding= 1),
        nn.ReLU()
    ) #14x14x16 to 14x14x8
    self.decoder_conv3 = nn.Sequential(
        nn.Conv2d(8,1, kernel_size = 3, stride = 1,padding= 1),
        nn.ReLU()
    ) #28x28x8 to 28x28x1

    #defining the unpooling operation
    self.unpool = nn.MaxUnpool2d(kernel_size = (2,2))

    def forward(self,x): #defines the forward pass and also the structure of
    →the network thus helping backprop

        encoded_input,indices1 = self.encoder_conv1(x.float()) # 28x28x1 to
    →14x14x8
        encoded_input,indices2 = self.encoder_conv2(encoded_input) #14x14x8 to
    →7x7x16
        encoded_input,indices3 = self.encoder_conv3(encoded_input) #7x7x16 to
    →3x3x16

        reconstructed_input = self.
    →unpool(encoded_input,indices3,output_size=torch.Size([batch_size, 16, 7,
    →7])) #3x3x16 to 7x7x16
        reconstructed_input = self.decoder_conv1(reconstructed_input)
    →#7x7x16 to 7x7x16
        reconstructed_input = self.unpool(reconstructed_input,indices2)
    →#7x7x16 to 14x14x16
        reconstructed_input = self.
    →decoder_conv2(reconstructed_input)#14x14x16 to 14x14x8
        reconstructed_input = self.
    →unpool(reconstructed_input,indices1)#14x14x8 to 28x28x8
        reconstructed_input = self.
    →decoder_conv3(reconstructed_input)#28x28x8 to 28x28x1


```


```
return reconstructed_input, encoded_input
```

```
[43]: from torchsummary import summary
print(conv_AE_with_unpooling().to(device))
```



```
conv_AE_with_unpooling(
  (encoder_conv1): Sequential(
    (0): Conv2d(1, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceil_mode=False)
  )
  (encoder_conv2): Sequential(
    (0): Conv2d(8, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceil_mode=False)
  )
  (encoder_conv3): Sequential(
    (0): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceil_mode=False)
  )
  (decoder_conv1): Sequential(
    (0): Identity()
  )
  (decoder_conv2): Sequential(
    (0): Conv2d(16, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
  )
  (decoder_conv3): Sequential(
    (0): Conv2d(8, 1, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
  )
  (unpool): MaxUnpool2d(kernel_size=(2, 2), stride=(2, 2), padding=(0, 0))
)
```

```
[45]: model_Q5_a = conv_AE_with_unpooling().to(device)
optimizer = torch.optim.Adam(model_Q5_a.parameters(), lr=learning_rate)
train_losses_AE_Q5_a , test_losses_AE_Q5_a =
↳train_test(model_Q5_a, device, train_loader, test_loader, optimizer, lossfn, q5_flag=True)
plot_losses(train_losses_AE_Q5_a, test_losses_AE_Q5_a, model_name =
↳"Convolutional AE with Unpooling")
```





```
[46]: index = random.randint(0,9999)
test_image = test_loader.dataset.data[index, :, :].clone()
plot_reconstructed_image(model_Q5_a,device,test_image,□
↪model_name="Convolutional AE with Unpooling",q5_flag=True)
```




```
[47]: class conv_AE_with_deconv(nn.Module):
    def __init__(self):
        super(conv_AE_with_deconv,self).__init__()

        #encoder
        self.encoder_conv1 = nn.Sequential(
            nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2))
        )
        self.encoder_conv2 = nn.Sequential(
            nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2))
        )
        self.encoder_conv3 = nn.Sequential(
            nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2))
        )

        #decoder module
        self.decoder_conv1 = nn.Sequential(
            nn.ConvTranspose2d(16,16, kernel_size = 3, stride = 2),
            nn.ReLU()
        )
        self.decoder_conv2 = nn.Sequential(
            nn.ConvTranspose2d(16,8, kernel_size = 4, stride = 2, padding = 1),
            nn.ReLU()
        )
        self.decoder_conv3 = nn.Sequential(
            nn.ConvTranspose2d(8,1, kernel_size = 4, stride = 2, padding = 1),
            nn.ReLU()
        )

    def forward(self,x):

        encoded_input  = self.encoder_conv1(x.float())
        encoded_input  = self.encoder_conv2(encoded_input)
        encoded_input  = self.encoder_conv3(encoded_input)


        reconstructed_input = self.decoder_conv1(encoded_input)
        reconstructed_input = self.decoder_conv2(reconstructed_input)
        reconstructed_input = self.decoder_conv3(reconstructed_input)


        return reconstructed_input,encoded_input
```

```
[48]: print(conv_AE_with_deconv().to(device))
```



```
conv_AE_with_deconv(
  (encoder_conv1): Sequential(
    (0): Conv2d(1, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceiling_mode=False)
  )
  (encoder_conv2): Sequential(
    (0): Conv2d(8, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceiling_mode=False)
  )
  (encoder_conv3): Sequential(
    (0): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceiling_mode=False)
  )
  (decoder_conv1): Sequential(
    (0): ConvTranspose2d(16, 16, kernel_size=(3, 3), stride=(2, 2))
    (1): ReLU()
  )
  (decoder_conv2): Sequential(
    (0): ConvTranspose2d(16, 8, kernel_size=(4, 4), stride=(2, 2), padding=(1,
1))
    (1): ReLU()
  )
  (decoder_conv3): Sequential(
    (0): ConvTranspose2d(8, 1, kernel_size=(4, 4), stride=(2, 2), padding=(1,
1))
    (1): ReLU()
  )
)
```

```
[49]: model_Q5_b = conv_AE_with_deconv().to(device)
optimizer = torch.optim.Adam(model_Q5_b.parameters(), lr=learning_rate)
train_losses_AE_Q5_b , test_losses_AE_Q5_b =
    ↪ train_test(model_Q5_b, device, train_loader, test_loader, optimizer, lossfn, q5_flag=True)
plot_losses(train_losses_AE_Q5_b, test_losses_AE_Q5_b, model_name =
    ↪ "Convolutional AE with deconvolution")
```





```
[50]: index = random.randint(0,9999)
test_image = test_loader.dataset.data[index, :, :].clone()
plot_reconstructed_image(model_Q5_b,device,test_image,□
↪model_name="Convolutional AE with deconvolution",q5_flag=True)
```



```
[51]: class conv_AE_with_deconv_unpool(nn.Module):
    def __init__(self):
        super(conv_AE_with_deconv_unpool,self).__init__()

        #encoder
        self.encoder_conv1 = nn.Sequential(
            nn.Conv2d(1,8, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        )
        self.encoder_conv2 = nn.Sequential(
            nn.Conv2d(8,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        )
        self.encoder_conv3 = nn.Sequential(
            nn.Conv2d(16,16, kernel_size = 3, stride = 1,padding= 1),
            nn.ReLU(),nn.MaxPool2d(kernel_size = (2,2),return_indices = True)
        )

        #initializing the decoder module
        self.decoder_conv1 = nn.Sequential(
            nn.ConvTranspose2d(16,16, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU()
        )
        self.decoder_conv2 = nn.Sequential(
            nn.ConvTranspose2d(16,8, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU()
        )
        self.decoder_conv3 = nn.Sequential(
            nn.ConvTranspose2d(8,1, kernel_size = 3, stride = 1, padding = 1),
            nn.ReLU()
        )

        #unpooling
        self.unpool = nn.MaxUnpool2d(kernel_size = (2,2))

        def forward(self,x): #defines the forward pass and also the structure of
        ↪the network thus helping backprop

            encoded_input,indices1 = self.encoder_conv1(x.float())
            encoded_input,indices2 = self.encoder_conv2(encoded_input)
            encoded_input,indices3 = self.encoder_conv3(encoded_input)

            reconstructed_input = self.
            ↪unpool(encoded_input,indices3,output_size=torch.Size([batch_size, 16, 7, 7]))
```

```

reconstructed_input = self.decoder_conv1(reconstructed_input)
reconstructed_input = self.unpool(reconstructed_input, indices2)
reconstructed_input = self.decoder_conv2(reconstructed_input)
reconstructed_input = self.unpool(reconstructed_input, indices1)
reconstructed_input = self.decoder_conv3(reconstructed_input)

return reconstructed_input, encoded_input

```

```

[52]: from torchsummary import summary
print(conv_AE_with_deconv_unpool().to(device))

```


```

conv_AE_with_deconv_unpool(
  (encoder_conv1): Sequential(
    (0): Conv2d(1, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceil_mode=False)
  )
  (encoder_conv2): Sequential(
    (0): Conv2d(8, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceil_mode=False)
  )
  (encoder_conv3): Sequential(
    (0): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceil_mode=False)
  )
  (decoder_conv1): Sequential(
    (0): ConvTranspose2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (1): ReLU()
  )
  (decoder_conv2): Sequential(
    (0): ConvTranspose2d(16, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (1): ReLU()
  )
  (decoder_conv3): Sequential(
    (0): ConvTranspose2d(8, 1, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (1): ReLU()
  )
  (unpool): MaxUnpool2d(kernel_size=(2, 2), stride=(2, 2), padding=(0, 0))
)


```

```
[53]: model_Q5_c = conv_AE_with_deconv_unpool().to(device)
optimizer = torch.optim.Adam(model_Q5_c.parameters(), lr=learning_rate)
train_losses_AE_Q5_c , test_losses_AE_Q5_c = □
↳train_test(model_Q5_c,device,train_loader,test_loader,optimizer,lossfn,q5_flag=True)
plot_losses(train_losses_AE_Q5_c, test_losses_AE_Q5_c, model_name =□
↳"Convolutional AE with deconvolution and unpooling")
```

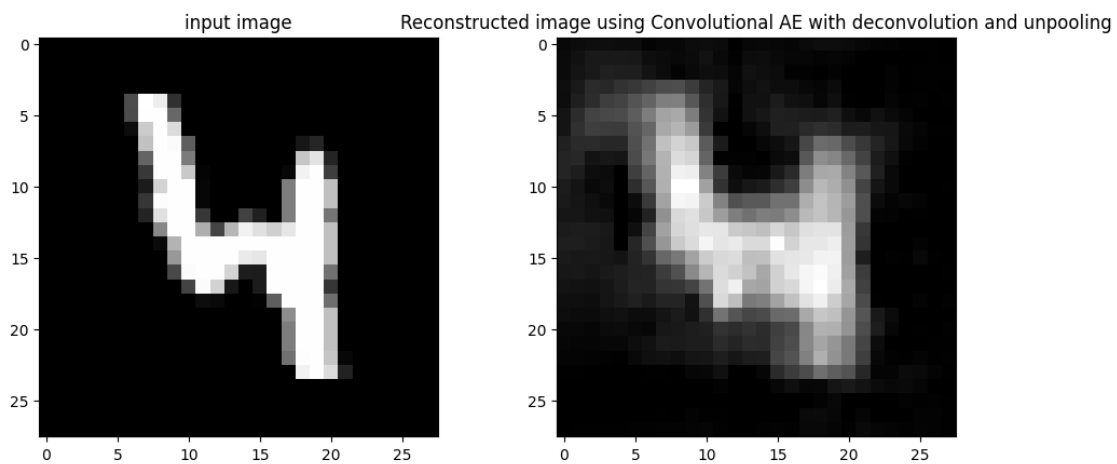
MSE train loss using Convolutional AE with deconvolution and unpooling



MSE test loss using Convolutional AE with deconvolution and unpooling



```
[54]: index = random.randint(0,9999)
test_image = test_loader.dataset.data[index, :, :].clone()
plot_reconstructed_image(model_Q5_c,device,test_image,□
↪model_name="Convolutional AE with deconvolution and unpooling",q5_flag=True)
```





```
[55]: from torchvision.utils import make_grid
```


```
[56]: def
    ↪conv_AE_visualize_decoder_weights(model,model_name,device,decoder_conv_filter,decoder_layer
    ↪
    if(device == torch.device('cuda')):
        decoder_conv_filter = decoder_conv_filter.cpu()
    #normalize the filter weights
    decoder_conv_filter -= decoder_conv_filter.min()
    decoder_conv_filter /= decoder_conv_filter.max()
    (x,y,z,w) = decoder_conv_filter.size()
    filt_ind = np.random.randint(0 ,decoder_conv_filter.size()[0],3)
    for ind in filt_ind:
        image = make_grid(decoder_conv_filter[ind].reshape(y,1,z,w))
        image = image.permute(1,2,0)
        plt.imshow(image)
        plt.title("Decoder "+str(decoder_layer)+" Convolutional layer filter_
    ↪outputs for filter no. "+str(ind)+" of "+str(model_name))
        plt.show()
```

```
[57]: conv_3_filter = model_Q5_a.decoder_conv3[0].weight.detach().clone()
conv_AE_visualize_decoder_weights(model=model_Q5_a,model_name="CAE with_
    ↪unpooling",device=device,decoder_conv_filter=conv_3_filter,decoder_layer="Third")
```


Decoder Third Convolutional layer filter outputs for filter no. 0 of CAE with unpooling



Decoder Third Convolutional layer filter outputs for filter no. 0 of CAE with unpooling




Decoder Third Convolutional layer filter outputs for filter no. 0 of CAE with unpooling




```
[58]: conv_1_filter = model_Q5_b.decoder_conv1[0].weight.detach().clone()
conv_AE_visualize_decoder_weights(model=model_Q5_b,model_name="CAE with_
↳deconvolution",device=device,decoder_conv_filter=conv_1_filter,decoder_layer="First")
```


Decoder First Convolutional layer filter outputs for filter no. 1 of CAE with deconvolution



Decoder First Convolutional layer filter outputs for filter no. 7 of CAE with deconvolution




Decoder First Convolutional layer filter outputs for filter no. 13 of CAE with deconvolution




```
[65]: conv_2_filter = model_Q5_b.decoder_conv2[0].weight.detach().clone()
conv_AE_visualize_decoder_weights(model=model_Q5_b,model_name="",device=device,decoder_conv_fi
```


Decoder Second Convolutional layer filter outputs for filter no. 8 of



Decoder Second Convolutional layer filter outputs for filter no. 2 of




Decoder Second Convolutional layer filter outputs for filter no. 3 of




```
[66]: conv_3_filter = model_Q5_b.decoder_conv3[0].weight.detach().clone()
conv_AE_visualize_decoder_weights(model=model_Q5_b,model_name="CAE with_
↳deconvolution",device=device,decoder_conv_filter=conv_3_filter,decoder_layer="Third")
```


Decoder Third Convolutional layer filter outputs for filter no. 1 of CAE with deconvolution



Decoder Third Convolutional layer filter outputs for filter no. 1 of CAE with deconvolution



Decoder Third Convolutional layer filter outputs for filter no. 7 of CAE with deconvolution



[]: