

Task_02

October 2, 2022

1 2 - Improving Performance

In the previous notebook, we got the fundamentals down for sentiment analysis. In this notebook, we'll actually get decent results.

We will use: - bidirectional RNN - multi-layer RNN

This will allow us to achieve ~84% test accuracy.

1.1 Preparing Data

```
[ ]: !pip install torchtext==0.10.0
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting torchtext==0.10.0

Downloading torchtext-0.10.0-cp37-cp37m-manylinux1_x86_64.whl (7.6 MB)

|| 7.6 MB 25.7 MB/s

Collecting torch==1.9.0

Downloading torch-1.9.0-cp37-cp37m-manylinux1_x86_64.whl (831.4 MB)

|| 831.4 MB 2.8 kB/s

Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from torchtext==0.10.0) (1.21.6)

Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from torchtext==0.10.0) (4.64.1)

Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from torchtext==0.10.0) (2.23.0)

Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from torch==1.9.0->torchtext==0.10.0) (4.1.1)

Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests->torchtext==0.10.0) (3.0.4)

Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests->torchtext==0.10.0) (1.24.3)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests->torchtext==0.10.0) (2022.6.15)

Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests->torchtext==0.10.0) (2.10)

Installing collected packages: torch, torchtext

Attempting uninstall: torch

Found existing installation: torch 1.12.1+cu113

Uninstalling torch-1.12.1+cu113:

Successfully uninstalled torch-1.12.1+cu113

Attempting uninstall: torchtext

Found existing installation: torchtext 0.13.1

Uninstalling torchtext-0.13.1:

Successfully uninstalled torchtext-0.13.1

ERROR: pip's dependency resolver does not currently take into account all

the packages that are installed. This behaviour is the source of the following dependency conflicts.

torchvision 0.13.1+cu113 requires torch==1.12.1, but you have torch 1.9.0 which is incompatible.

torchaudio 0.12.1+cu113 requires torch==1.12.1, but you have torch 1.9.0 which is incompatible.

Successfully installed torch-1.9.0 torchtext-0.10.0

```
[ ]: import torch
from torchtext.legacy import data

SEED = 1234

torch.manual_seed(SEED)
torch.backends.cudnn.deterministic = True

TEXT = data.Field(tokenize = 'spacy',
                  tokenizer_language = 'en_core_web_sm')

LABEL = data.LabelField(dtype = torch.float)
```

```
[ ]: from torchtext.legacy import datasets

train_data, test_data = datasets.IMDB.splits(TEXT, LABEL)

downloading aclImdb_v1.tar.gz
aclImdb_v1.tar.gz: 100%|| 84.1M/84.1M [00:03<00:00, 23.4MB/s]
```

```
[ ]: import random

train_data, valid_data = train_data.split(random_state = random.seed(SEED))
```

Next is the use of pre-trained word embeddings. Now, instead of having our word embeddings

initialized randomly, they are initialized with these pre-trained vectors. We get these vectors simply by specifying which vectors we want and passing it as an argument to `build_vocab`. `TorchText` handles downloading the vectors and associating them with the correct words in our vocabulary.

Here, we'll be using the "glove.6B.100d" vectors". `glove` is the algorithm used to calculate the vectors, go [here](#) for more. `6B` indicates these vectors were trained on 6 billion tokens and `100d` indicates these vectors are 100-dimensional.

You can see the other available vectors [here](#).

The theory is that these pre-trained vectors already have words with similar semantic meaning close together in vector space, e.g. "terrible", "awful", "dreadful" are nearby. This gives our embedding layer a good initialization as it does not have to learn these relations from scratch.

```
[ ]: MAX_VOCAB_SIZE = 25_000

TEXT.build_vocab(train_data,
                  max_size = MAX_VOCAB_SIZE,
                  vectors = "glove.6B.100d",
                  unk_init = torch.Tensor.normal_)

LABEL.build_vocab(train_data)
```

```
.vector_cache/glove.6B.zip: 862MB [02:43, 5.28MB/s]
100%|| 399999/400000 [00:13<00:00, 29965.99it/s]
```

```
[ ]: BATCH_SIZE = 64

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

train_iterator, valid_iterator, test_iterator = data.BucketIterator.splits(
    (train_data, valid_data, test_data),
    batch_size = BATCH_SIZE,
    sort_within_batch = True,
    device = device)
```

1.2 Build the Model

1.2.1 Different RNN Architecture

We'll be using a different RNN architecture called a Long Short-Term Memory (LSTM). Why is an LSTM better than a standard RNN? Standard RNNs suffer from the [vanishing gradient problem](#). LSTMs overcome this by having an extra recurrent state called a *cell*, c - which can be thought of as the "memory" of the LSTM - and they use multiple *gates* which control the flow of information into and out of the memory. For more information, go [here](#). We can simply think of the LSTM as a function of x_t , h_t and c_t , instead of just x_t and h_t .

$$(h_t, c_t) = \text{LSTM}(x_t, h_t, c_t)$$

Thus, the model using an LSTM looks something like (with the embedding layers omitted):

The initial cell state, c_0 , like the initial hidden state is initialized to a tensor of all zeros. The sentiment prediction is still, however, only made using the final hidden state, not the final cell state, i.e. $\hat{y} = f(h_T)$.

1.2.2 Bidirectional RNN

The concept behind a bidirectional RNN is simple. As well as having an RNN processing the words in the sentence from the first to the last (a forward RNN), we have a second RNN processing the words in the sentence from the **last to the first** (a backward RNN). At time step t , the forward RNN is processing word x_t , and the backward RNN is processing word x_{T-t+1} .

In PyTorch, the hidden state (and cell state) tensors returned by the forward and backward RNNs are stacked on top of each other in a single tensor.

We make our sentiment prediction using a concatenation of the last hidden state from the forward RNN (obtained from final word of the sentence), h_T^{\rightarrow} , and the last hidden state from the backward RNN (obtained from the first word of the sentence), h_T^{\leftarrow} , i.e. $\hat{y} = f(h_T^{\rightarrow}, h_T^{\leftarrow})$

The image below shows a bi-directional RNN, with the forward RNN in orange, the backward RNN in green and the linear layer in silver.

1.2.3 Multi-layer RNN

Multi-layer RNNs (also called *deep RNNs*) are another simple concept. The idea is that we add additional RNNs on top of the initial standard RNN, where each RNN added is another *layer*. The hidden state output by the first (bottom) RNN at time-step t will be the input to the RNN above it at time step t . The prediction is then made from the final hidden state of the final (highest) layer.

The image below shows a multi-layer unidirectional RNN, where the layer number is given as a superscript. Also note that each layer needs their own initial hidden state, h_0^L .

```
[ ]: import torch.nn as nn

class RNN(nn.Module):
    def __init__(self, vocab_size, embedding_dim, hidden_dim, output_dim,
        ↪ n_layers,
            bidirectional, dropout, pad_idx):

        super().__init__()

        self.embedding = nn.Embedding(vocab_size, embedding_dim, padding_idx =
        ↪ pad_idx) ### CODE HERE ###

        ### CODE HERE ###

        self.rnn = nn.LSTM(embedding_dim,
```

```

        hidden_dim,
        num_layers=n_layers,
        bidirectional=bidirectional,
        dropout=dropout)

    self.fc = nn.Linear(hidden_dim * 2, output_dim) ### CODE HERE ###

    self.dropout = nn.Dropout(dropout)

    def forward(self, text):

        #text = [sent len, batch size]

        embedded = self.dropout(self.embedding(text))

        #embedded = [sent len, batch size, emb dim]

        output, (hidden, cell) = self.rnn(embedded) ### CODE HERE ###

        #output = [sent len, batch size, hid dim * num directions]
        #output over padding tokens are zero tensors

        #hidden = [num layers * num directions, batch size, hid dim]
        #cell = [num layers * num directions, batch size, hid dim]

        #concat the final forward (hidden[-2,:,:]) and backward (hidden[-1,:,:])
        → hidden layers
        #and apply dropout

        hidden = self.dropout(torch.cat((hidden[-2,:,:], hidden[-1,:,:]), dim =
        → 1))

        #hidden = [batch size, hid dim * num directions]

        return self.fc(hidden)

```

```

[ ]: INPUT_DIM = len(TEXT.vocab)
    EMBEDDING_DIM = 100
    HIDDEN_DIM = 256
    OUTPUT_DIM = 1
    N_LAYERS = 2
    BIDIRECTIONAL = True
    DROPOUT = 0.5
    PAD_IDX = TEXT.vocab.stoi[TEXT.pad_token]

    model = RNN(INPUT_DIM,
                 EMBEDDING_DIM,

```

```
HIDDEN_DIM,
OUTPUT_DIM,
N_LAYERS,
BIDIRECTIONAL,
DROPOUT,
PAD_IDX)
```

We'll print out the number of parameters in our model.

Notice how we have almost twice as many parameters as before!

```
[ ]: def count_parameters(model):
      return sum(p.numel() for p in model.parameters() if p.requires_grad)

      print(f'The model has {count_parameters(model):,} trainable parameters')
```

The model has 4,810,857 trainable parameters

The final addition is copying the pre-trained word embeddings we loaded earlier into the `embedding` layer of our model.

We retrieve the embeddings from the field's vocab, and check they're the correct size, *[vocab size, embedding dim]*

```
[ ]: pretrained_embeddings = TEXT.vocab.vectors

      print(pretrained_embeddings.shape)
```

```
torch.Size([25002, 100])
```

We then replace the initial weights of the `embedding` layer with the pre-trained embeddings.

Note: this should always be done on the `weight.data` and not the `weight`!

```
[ ]: model.embedding.weight.data.copy_(pretrained_embeddings)

[ ]: tensor([[ -0.1117, -0.4966,  0.1631, ...,  1.2647, -0.2753, -0.1325],
            [-0.8555, -0.7208,  1.3755, ...,  0.0825, -1.1314,  0.3997],
            [-0.0382, -0.2449,  0.7281, ..., -0.1459,  0.8278,  0.2706],
            ...,
            [-0.0614, -0.0516, -0.6159, ..., -0.0354,  0.0379, -0.1809],
            [-0.4127,  0.0867, -0.2232, ..., -0.2191,  0.0485,  1.2073],
            [-0.1182, -0.4701, -0.0600, ...,  0.7991, -0.0194,  0.4785]])
```

```
[ ]: UNK_IDX = TEXT.vocab.stoi[TEXT.unk_token]

      model.embedding.weight.data[UNK_IDX] = torch.zeros(EMBEDDING_DIM)
      model.embedding.weight.data[PAD_IDX] = torch.zeros(EMBEDDING_DIM)

      print(model.embedding.weight.data)
```

```
tensor([[ 0.0000,  0.0000,  0.0000, ...,  0.0000,  0.0000,  0.0000],
        [ 0.0000,  0.0000,  0.0000, ...,  0.0000,  0.0000,  0.0000],
        [-0.0382, -0.2449,  0.7281, ..., -0.1459,  0.8278,  0.2706],
        ...,
        [-0.0614, -0.0516, -0.6159, ..., -0.0354,  0.0379, -0.1809],
        [-0.4127,  0.0867, -0.2232, ..., -0.2191,  0.0485,  1.2073],
        [-0.1182, -0.4701, -0.0600, ...,  0.7991, -0.0194,  0.4785]])
```

1.3 Train the Model

```
[ ]: import torch.optim as optim
```

```
optimizer = optim.SGD(model.parameters(), lr=1e-3)
```

```
[ ]: criterion = nn.BCEWithLogitsLoss()
```

```
model = model.to(device)
```

```
criterion = criterion.to(device)
```

```
[ ]: def binary_accuracy(preds, y):
```

```
    #round predictions to the closest integer
```

```
    rounded_preds = torch.round(torch.sigmoid(preds))
```

```
    correct = (rounded_preds == y).float() #convert into float for division
```

```
    acc = correct.sum() / len(correct)
```

```
    return acc
```

```
[ ]: from tqdm import tqdm
```

```
[ ]: def train(model, iterator, optimizer, criterion):
```

```
    epoch_loss = 0
```

```
    epoch_acc = 0
```

```
    model.train()
```

```
    for batch in tqdm(iterator):
```

```
        optimizer.zero_grad()
```

```
        predictions = model(batch.text).squeeze(1)
```

```
        loss = criterion(predictions, batch.label)
```

```
        acc = binary_accuracy(predictions, batch.label)
```

```
        loss.backward()
```

```

optimizer.step()

epoch_loss += loss.item()
epoch_acc += acc.item()

return epoch_loss / len(iterator), epoch_acc / len(iterator)

```

```

[ ]: def evaluate(model, iterator, criterion):

    epoch_loss = 0
    epoch_acc = 0

    model.eval()

    with torch.no_grad():

        for batch in tqdm(iterator):

            predictions = model(batch.text).squeeze(1)

            loss = criterion(predictions, batch.label)

            acc = binary_accuracy(predictions, batch.label)

            epoch_loss += loss.item()
            epoch_acc += acc.item()

    return epoch_loss / len(iterator), epoch_acc / len(iterator)

```

```

[ ]: import time

def epoch_time(start_time, end_time):
    elapsed_time = end_time - start_time
    elapsed_mins = int(elapsed_time / 60)
    elapsed_secs = int(elapsed_time - (elapsed_mins * 60))
    return elapsed_mins, elapsed_secs

```

```

[ ]: N_EPOCHS = 5

best_valid_loss = float('inf')

for epoch in range(N_EPOCHS):

    start_time = time.time()

    train_loss, train_acc = train(model, train_iterator, optimizer, criterion)
    valid_loss, valid_acc = evaluate(model, valid_iterator, criterion)

```



```

end_time = time.time()

epoch_mins, epoch_secs = epoch_time(start_time, end_time)

if valid_loss < best_valid_loss:
    best_valid_loss = valid_loss
    torch.save(model.state_dict(), 'tut7-model.pt')

print(f'Epoch: {epoch+1:02} | Epoch Time: {epoch_mins}m {epoch_secs}s')
print(f'\tTrain Loss: {train_loss:.3f} | Train Acc: {train_acc*100:.2f}%')
print(f'\tVal. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*100:.2f}%')

```

100%| 274/274 [00:33<00:00, 8.15it/s]

100%| 118/118 [00:04<00:00, 24.61it/s]

Epoch: 01 | Epoch Time: 0m 38s

Train Loss: 0.693 | Train Acc: 50.38%

Val. Loss: 0.693 | Val. Acc: 51.65%

100%| 274/274 [00:33<00:00, 8.27it/s]

100%| 118/118 [00:04<00:00, 26.07it/s]

Epoch: 02 | Epoch Time: 0m 37s

Train Loss: 0.693 | Train Acc: 49.60%

Val. Loss: 0.693 | Val. Acc: 51.95%

100%| 274/274 [00:33<00:00, 8.27it/s]

100%| 118/118 [00:04<00:00, 25.25it/s]

Epoch: 03 | Epoch Time: 0m 37s

Train Loss: 0.693 | Train Acc: 50.32%

Val. Loss: 0.693 | Val. Acc: 51.47%

100%| 274/274 [00:33<00:00, 8.22it/s]

100%| 118/118 [00:04<00:00, 25.48it/s]

Epoch: 04 | Epoch Time: 0m 37s

Train Loss: 0.693 | Train Acc: 50.57%

Val. Loss: 0.693 | Val. Acc: 50.92%

100%| 274/274 [00:33<00:00, 8.28it/s]

100%| 118/118 [00:04<00:00, 24.86it/s]

Epoch: 05 | Epoch Time: 0m 37s

Train Loss: 0.693 | Train Acc: 49.54%

Val. Loss: 0.693 | Val. Acc: 51.06%

```
[ ]: model.load_state_dict(torch.load('tut7-model.pt'))
```

```
test_loss, test_acc = evaluate(model, test_iterator, criterion)

print(f'Test Loss: {test_loss:.3f} | Test Acc: {test_acc*100:.2f}%')
```

100%|| 391/391 [00:14<00:00, 26.50it/s]

Test Loss: 0.693 | Test Acc: 49.44%