

EE21S055_Tutorial6_Bonus

September 18, 2022

0.1 Imports

```
[1]: import torch
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import Dataset, DataLoader
import torch.nn as nn
import torch.nn.functional as F
import sys
import numpy as np
import os
```

0.2 Utilising GPU using Pytorch

```
[2]: # cpu-gpu
a = torch.randn((3, 4))
print(a.device)

device = torch.device("cuda")
a = a.to(device)
print(a.device)

# a more generic code
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

cpu
cuda:0

```
[3]: torch.cuda.is_available()
```

[3]: True

```
[4]: !nvidia-smi
```

Sun Sep 18 11:55:15 2022

```
+-----+
| NVIDIA-SMI 460.32.03      Driver Version: 460.32.03      CUDA Version: 11.2      |
+-----+-----+-----+
```

GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile Uncorr. ECC
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util Compute M.
					MIG M.
=====					
0	Tesla T4	Off	00000000:00:04.0	Off	0
N/A	45C	P0	26W / 70W	612MiB / 15109MiB	1% Default
					N/A

Processes:					
GPU	GI	CI	PID	Type	Process name
	ID	ID			
=====					

GPU Memory Usage					

0.3 Dataset and Transforms

```
[5]: train_transform = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])
test_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

train_dset = torchvision.datasets.CIFAR10(root="data/", train=True,
    ↪transform=train_transform, download=True)
test_dset = torchvision.datasets.CIFAR10(root="data/", train=False,
    ↪transform=test_transform, download=True)
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to data/cifar-10-python.tar.gz

```
0%|          | 0/170498071 [00:00<?, ?it/s]
```

Extracting data/cifar-10-python.tar.gz to data/
Files already downloaded and verified

```
[6]: print(f"# of train samples: {len(train_dset)}")
    print(f"# of test samples: {len(test_dset)}")
```

```
# of train samples: 50000
# of test samples: 10000
```

```
[7]: train_loader = DataLoader(train_dset, batch_size=100, shuffle=True,
    ↪ num_workers=2)
test_loader = DataLoader(test_dset, batch_size=100, shuffle=False,
    ↪ num_workers=2)
```

```
[8]: print(f"# of train batches: {len(train_loader)}")
print(f"# of test batches: {len(test_loader)}")
```

```
# of train batches: 500
# of test batches: 100
```

```
[9]: print("sample i/o sizes")
data = next(iter(train_loader))
img, target = data
print(f"input size: {img.shape}")
print(f"output size: {target.shape}")
```

```
sample i/o sizes
input size: torch.Size([100, 3, 32, 32])
output size: torch.Size([100])
```

0.4 LeNet

```
[43]: class LeNet(nn.Module):
    def __init__(self):
        super(LeNet, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, kernel_size=5)
        self.conv2 = nn.Conv2d(6, 32, kernel_size=5)
        self.conv3 = nn.Conv2d(32, 64, kernel_size=5)
        # TODO: missing input feature size
        self.fc1 = nn.Linear(64*5*5, 120)
        self.fc2 = nn.Linear(120, 84)
        # TODO: missing output feature size
        self.fc3 = nn.Linear(84, 10)
        self.activ = nn.ReLU()

        # TODO: add maxpool operation of given kernel size
        # https://pytorch.org/docs/stable/nn.functional.html
    def pool(self, x, kernel_size=2):
        out = F.max_pool2d(x, kernel_size)
        return out

    def forward(self, x):
        out = self.activ(self.conv1(x))
        out = self.pool(out)
        out = self.activ(self.conv2(out))
        out = self.pool(out)
```

```

out = self.activ(self.conv3(out))
out = self.pool(out)

# TODO: flatten
out = out.view(out.size(0),-1)
out = self.activ(self.fc1(out))
out = self.activ(self.fc2(out))
out = self.fc3(out)
return out

```

0.5 Utility functions (can ignore)

```

[20]: def pbar(p=0, msg="", bar_len=20):
    sys.stdout.write("\033[K")
    sys.stdout.write("\x1b[2K" + "\r")
    block = int(round(bar_len * p))
    text = "Progress: [{}] {}% {}".format(
        "\x1b[32m" + "=" * (block - 1) + ">" + "\033[0m" + "-" * (bar_len -
↪block),
        round(p * 100, 2),
        msg,
    )
    print(text, end="\r")
    if p == 1:
        print()

class AvgMeter:
    def __init__(self):
        self.reset()

    def reset(self):
        self.metrics = {}

    def add(self, batch_metrics):
        if self.metrics == {}:
            for key, value in batch_metrics.items():
                self.metrics[key] = [value]
        else:
            for key, value in batch_metrics.items():
                self.metrics[key].append(value)

    def get(self):
        return {key: np.mean(value) for key, value in self.metrics.items()}

```

```

def msg(self):
    avg_metrics = {key: np.mean(value) for key, value in self.metrics.
↪items()}
    return "".join(["[{}] {:.5f} ".format(key, value) for key, value in
↪avg_metrics.items()])

```

0.6 Training

```

[21]: def train(model, optim, lr_sched=None, epochs=20, device=torch.device("cuda" if
↪torch.cuda.is_available() else "cpu"), criterion=None, metric_meter=None,
↪out_dir="out/"):
    model.to(device)
    best_acc = 0
    for epoch in range(epochs):
        model.train()
        metric_meter.reset()
        for indx, (img, target) in enumerate(train_loader):
            # TODO: send to device (cpu or gpu)
            img = img.to(device)
            target = target.to(device)

            # TODO: missing forward pass
            out = model(img)
            loss = criterion(out, target)
            # TODO: missing backward, parameter update
            optim.zero_grad()
            loss.backward()
            optim.step()
            metric_meter.add({"train loss": loss.item()})
            pbar(indx / len(train_loader), msg=metric_meter.msg())
        pbar(1, msg=metric_meter.msg())

    model.eval()
    metric_meter.reset()
    for indx, (img, target) in enumerate(test_loader):
        # TODO: send to device (cpu or gpu)
        img = img.to(device)
        target = target.to(device)

        # TODO: missing forward pass
        out = model(img)
        loss = criterion(out, target)
        # TODO: compute accuracy
        acc = (out.argmax(1) == target).type(torch.float).sum().item()

        metric_meter.add({"test loss": loss.item(), "test acc": acc})

```

```

        pbar(indx / len(test_loader), msg=metric_meter.msg())
    pbar(1, msg=metric_meter.msg())

    test_metrics = metric_meter.get()
    if test_metrics["test acc"] > best_acc:
        print(
            "\x1b[33m"
            + f"test acc improved from {round(best_acc, 5)} to \
↳{round(test_metrics['test acc'], 5)}"
            + "\033[0m"
        )
        best_acc = test_metrics['test acc']
        torch.save(model.state_dict(), os.path.join(out_dir, "best.ckpt"))
        lr_sched.step()

```

0.7 Run Experiments

```

[22]: def run_experiment(model_name="lenet", model_cfg=None, epochs=20):
        if model_name == "lenet":
            model = LeNet()
            optim = torch.optim.SGD(model.parameters(), lr=1e-1, momentum=0.9, \
↳weight_decay=5e-4)
            lr_sched = torch.optim.lr_scheduler.CosineAnnealingLR(optim, T_max=epochs)
            criterion = nn.CrossEntropyLoss()
            metric_meter = AvgMeter()
            out_dir = f"{model_name}_{model_cfg}"
            os.makedirs(out_dir, exist_ok=True)
            train(model, optim, lr_sched, epochs=epochs, criterion=criterion, \
↳metric_meter=metric_meter, out_dir=out_dir)

```

```

[23]: run_experiment(model_name="lenet")

```

```

Progress: [=====>] 100% [train loss] 2.09773
Progress: [=====>] 100% [test loss] 2.00331 [test acc]
25.95000
test acc improved from 0 to 25.95
Progress: [=====>] 100% [train loss] 1.90805
Progress: [=====>] 100% [test loss] 1.88315 [test acc]
29.02000
test acc improved from 25.95 to 29.02
Progress: [=====>] 100% [train loss] 1.87238
Progress: [=====>] 100% [test loss] 1.82579 [test acc]
33.49000
test acc improved from 29.02 to 33.49
Progress: [=====>] 100% [train loss] 1.84685
Progress: [=====>] 100% [test loss] 1.84448 [test acc]
31.75000

```

```

Progress: [=====>] 100% [train loss] 1.81004
Progress: [=====>] 100% [test loss] 1.71023 [test acc]
36.34000
test acc improved from 33.49 to 36.34
Progress: [=====>] 100% [train loss] 1.76901
Progress: [=====>] 100% [test loss] 1.74466 [test acc]
36.95000
test acc improved from 36.34 to 36.95
Progress: [=====>] 100% [train loss] 1.72513
Progress: [=====>] 100% [test loss] 1.65723 [test acc]
37.38000
test acc improved from 36.95 to 37.38
Progress: [=====>] 100% [train loss] 1.68298
Progress: [=====>] 100% [test loss] 1.55102 [test acc]
43.42000
test acc improved from 37.38 to 43.42
Progress: [=====>] 100% [train loss] 1.62786
Progress: [=====>] 100% [test loss] 1.54186 [test acc]
44.17000
test acc improved from 43.42 to 44.17
Progress: [=====>] 100% [train loss] 1.57178
Progress: [=====>] 100% [test loss] 1.54262 [test acc]
44.66000
test acc improved from 44.17 to 44.66
Progress: [=====>] 100% [train loss] 1.52041
Progress: [=====>] 100% [test loss] 1.47860 [test acc]
46.27000
test acc improved from 44.66 to 46.27
Progress: [=====>] 100% [train loss] 1.46437
Progress: [=====>] 100% [test loss] 1.41450 [test acc]
48.96000
test acc improved from 46.27 to 48.96
Progress: [=====>] 100% [train loss] 1.39063
Progress: [=====>] 100% [test loss] 1.31499 [test acc]
53.62000
test acc improved from 48.96 to 53.62
Progress: [=====>] 100% [train loss] 1.31890
Progress: [=====>] 100% [test loss] 1.21614 [test acc]
56.95000
test acc improved from 53.62 to 56.95
Progress: [=====>] 100% [train loss] 1.27173
Progress: [=====>] 100% [test loss] 1.18751 [test acc]
58.10000
test acc improved from 56.95 to 58.1
Progress: [=====>] 100% [train loss] 1.20838
Progress: [=====>] 100% [test loss] 1.14827 [test acc]
59.24000
test acc improved from 58.1 to 59.24

```

```

Progress: [=====>] 100% [train loss] 1.15214
Progress: [=====>] 100% [test loss] 1.10608 [test acc]
61.27000
test acc improved from 59.24 to 61.27
Progress: [=====>] 100% [train loss] 1.11630
Progress: [=====>] 100% [test loss] 1.04973 [test acc]
62.81000
test acc improved from 61.27 to 62.81
Progress: [=====>] 100% [train loss] 1.07916
Progress: [=====>] 100% [test loss] 1.03096 [test acc]
63.75000
test acc improved from 62.81 to 63.75
Progress: [=====>] 100% [train loss] 1.06248
Progress: [=====>] 100% [test loss] 1.02110 [test acc]
64.48000
test acc improved from 63.75 to 64.48

16->32 Accuracy Moved from 62.57% to 64.48%

```

```

[37]: def run_experiment(model_name="lenet", model_cfg=None, epochs=20):
        if model_name == "lenet":
            model = LeNet()
            optim = torch.optim.SGD(model.parameters(), lr=1e-1, momentum=0.9,
            ↪weight_decay=5e-4)
            lr_sched = torch.optim.lr_scheduler.CosineAnnealingLR(optim, T_max=epochs)
            criterion = nn.CrossEntropyLoss()
            metric_meter = AvgMeter()
            out_dir = f"{model_name}_{model_cfg}"
            os.makedirs(out_dir, exist_ok=True)
            train(model, optim, lr_sched, epochs=epochs, criterion=criterion,
            ↪metric_meter=metric_meter, out_dir=out_dir)

```

```

[42]: run_experiment(model_name="lenet")

```

```

Progress: [=====>] 100% [train loss] 2.00005
Progress: [=====>] 100% [test loss] 1.84241 [test acc]
27.50000
test acc improved from 0 to 27.5
Progress: [=====>] 100% [train loss] 1.88387
Progress: [=====>] 100% [test loss] 1.87203 [test acc]
28.54000
test acc improved from 27.5 to 28.54
Progress: [=====>] 100% [train loss] 1.83851
Progress: [=====>] 100% [test loss] 1.86013 [test acc]
32.68000
test acc improved from 28.54 to 32.68
Progress: [=====>] 100% [train loss] 1.75704
Progress: [=====>] 100% [test loss] 1.65365 [test acc]

```



```

40.93000
test acc improved from 32.68 to 40.93
Progress: [=====>] 100% [train loss] 1.69457
Progress: [=====>] 100% [test loss] 1.64754 [test acc]
41.67000
test acc improved from 40.93 to 41.67
Progress: [=====>] 100% [train loss] 1.65990
Progress: [=====>] 100% [test loss] 1.60931 [test acc]
43.41000
test acc improved from 41.67 to 43.41
Progress: [=====>] 100% [train loss] 1.59530
Progress: [=====>] 100% [test loss] 1.54697 [test acc]
44.27000
test acc improved from 43.41 to 44.27
Progress: [=====>] 100% [train loss] 1.54157
Progress: [=====>] 100% [test loss] 1.45000 [test acc]
46.65000
test acc improved from 44.27 to 46.65
Progress: [=====>] 100% [train loss] 1.48376
Progress: [=====>] 100% [test loss] 1.45143 [test acc]
49.88000
test acc improved from 46.65 to 49.88
Progress: [=====>] 100% [train loss] 1.43155
Progress: [=====>] 100% [test loss] 1.33434 [test acc]
52.40000
test acc improved from 49.88 to 52.4
Progress: [=====>] 100% [train loss] 1.36831
Progress: [=====>] 100% [test loss] 1.26346 [test acc]
56.50000
test acc improved from 52.4 to 56.5
Progress: [=====>] 100% [train loss] 1.29312
Progress: [=====>] 100% [test loss] 1.16919 [test acc]
59.38000
test acc improved from 56.5 to 59.38
Progress: [=====>] 100% [train loss] 1.21105
Progress: [=====>] 100% [test loss] 1.08434 [test acc]
63.02000
test acc improved from 59.38 to 63.02
Progress: [=====>] 100% [train loss] 1.14628
Progress: [=====>] 100% [test loss] 1.07678 [test acc]
62.69000
Progress: [=====>] 100% [train loss] 1.07332
Progress: [=====>] 100% [test loss] 1.00292 [test acc]
65.20000
test acc improved from 63.02 to 65.2
Progress: [=====>] 100% [train loss] 1.01084
Progress: [=====>] 100% [test loss] 0.93889 [test acc]
67.80000

```

```
test acc improved from 65.2 to 67.8
Progress: [=====>] 100% [train loss] 0.95182
Progress: [=====>] 100% [test loss] 0.88403 [test acc]
69.45000
test acc improved from 67.8 to 69.45
Progress: [=====>] 100% [train loss] 0.91050
Progress: [=====>] 100% [test loss] 0.85413 [test acc]
70.31000
test acc improved from 69.45 to 70.31
Progress: [=====>] 100% [train loss] 0.86950
Progress: [=====>] 100% [test loss] 0.83311 [test acc]
70.96000
test acc improved from 70.31 to 70.96
Progress: [=====>] 100% [train loss] 0.85926
Progress: [=====>] 100% [test loss] 0.82271 [test acc]
71.33000
test acc improved from 70.96 to 71.33

32->64 Accuracy Moved from 64.48% to 71.33
```

```
[ ]: 
```

```
[ ]: 
```